


Per-study stage latency — GPU inference scales with study size, CPU preprocess stays ~flat

M2 Pro / mps



GB10 / cuda

